

# INTRODUÇÃO À PROGRAMAÇÃO PARA GPUs USANDO CUDA

---

Pedro Bruel

phrb@ime.usp.br

29 de Setembro de 2015



Instituto de Matemática e Estatística  
Universidade de São Paulo

## 1. Introdução

- *Template* para programas CUDA
- *Profilers* e *Debuggers*

## 1. Introdução

- *Template* para programas CUDA
- *Profilers* e *Debuggers*

## 2. Ferramentas

- `nvcc`
- `cuda-gdb`
- `cuda-memcheck`

Os *pdfs* com as aulas e todo o código fonte usado nos exemplos estão no [GitHub](#):

- [github.com/phrb/aulas-gpu](https://github.com/phrb/aulas-gpu)

Os *pdfs* com as aulas e todo o código fonte usado nos exemplos estão no [GitHub](#):

- [github.com/phrb/aulas-gpu](https://github.com/phrb/aulas-gpu)

Outros recursos:

- CUDA Toolkit Documentation: [docs.nvidia.com/cuda](https://docs.nvidia.com/cuda)
- GPU Teaching Kit: [syllabus.gputeachingkit.com](https://syllabus.gputeachingkit.com)
- iPython: [ipython.org/notebook.html](https://ipython.org/notebook.html)
- CUDA Toolkit: [developer.nvidia.com/cuda-toolkit](https://developer.nvidia.com/cuda-toolkit)
- Anaconda: [continuum.io/downloads](https://continuum.io/downloads)

Os próximos *slides* foram adaptados do material disponível no GPU Teaching Kit:

- [syllabus.gputeachingkit.com](https://syllabus.gputeachingkit.com)



## GPU Teaching Kit



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).



GPU Teaching Kit

Accelerated Computing



# Lecture 2.4 – Introduction to CUDA C

Introduction to the CUDA Toolkit



# Objective

- To become familiar with some valuable tools and resources from the CUDA Toolkit
  - Compiler flags
  - Debuggers
  - Profilers

# GPU Programming Languages

---

Numerical analytics ►

MATLAB, Mathematica, LabVIEW

Fortran ►

CUDA Fortran

C ►

CUDA C

C++ ►

CUDA C++

Python ►

PyCUDA, Copperhead, Numba, NumbaPro

F# ►

Alea.cuBase

# CUDA - C

## Applications

Libraries

Easy to use  
Most Performance

Compiler  
Directives

Easy to use  
Portable code

Programming  
Languages

Most Performance  
Most Flexibility

# NVCC Compiler

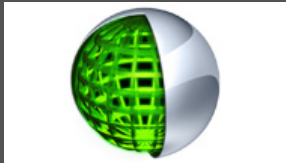
- NVIDIA provides a CUDA-C compiler
  - `nvcc`
- NVCC compiles device code then forwards code on to the host compiler (e.g. `g++`)
- Can be used to compile & link host only applications

# Compiler Flags

- Remember there are two compilers being used
  - NVCC: Device code
  - Host Compiler: C/C++ code
- NVCC supports some host compiler flags
  - If flag is unsupported, use `-Xcompiler` to forward to host
    - e.g. `-Xcompiler -fopenmp`
- Debugging Flags
  - `-g`: Include host debugging symbols
  - `-G`: Include device debugging symbols
  - `-lineinfo`: Include line information with symbols

# Developer Tools - Debuggers

**NSIGHT**



**CUDA-GDB**



**CUDA MEMCHECK**



**NVIDIA Provided**

**allinea**  
DDT

**TotalView®**

**3<sup>rd</sup> Party**

<https://developer.nvidia.com/debugging-solutions>

# CUDA-MEMCHECK

- Memory debugging tool
  - No recompilation necessary
    - `%> cuda-memcheck ./exe`
- Can detect the following errors
  - Memory leaks
  - Memory errors (OOB, misaligned access, illegal instruction, etc)
  - Race conditions
  - Illegal Barriers
  - Uninitialized Memory
- For line numbers use the following compiler flags:
  - `-Xcompiler -rdynamic -lineinfo`

<http://docs.nvidia.com/cuda/cuda-memcheck>

# CUDA-GDB

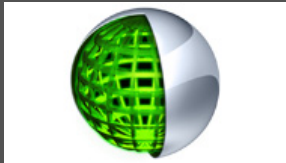
- cuda-gdb is an extension of GDB
  - Provides seamless debugging of CUDA and CPU code
- Works on Linux and Macintosh
  - For a Windows debugger use NSIGHT Visual Studio Edition

<http://docs.nvidia.com/cuda/cuda-gdb>

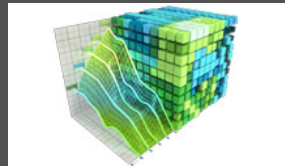


# Developer Tools - Profilers

**NSIGHT**



**NVVP**

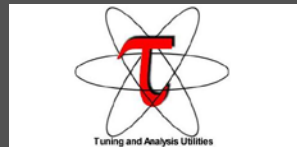


**NVPROF**

```
--20561== Profiling result:
Time(s)    Time    Calls    Avg    Min    Max    Name
49.88%    866.6ms    504758    1.727ms    1.560ms    2.815ms    void th
unt_thrust::detail::device_generate_function::thrust::detail::fill
25.33%    449.0ms    252662    1.741ms    1.536ms    2.308ms    void th
t_thrust::detail::device_generate_function::thrust::detail::fill_fo
17.87%    296.6ms    200    1.483ms    1.264ms    1.725ms    kerComp
2.98%    51.81ms    200    259.89us    246.97us    264.83us    kerTake
1.16%    20.17ms    501    40.26us    30us    17.47ms    [CUDA m
0.93%    16.19ms    200    80.99us    71.84us    90.75us    kerColl
0.73%    12.63ms    400    31.58us    14.72us    50.43us    [CUDA m
0.69%    12.07ms    200    60.37us    59.68us    62.36us    kerReme
0.63%    10.99ms    200    54.96us    52.68us    58.20us    kerTake
0.32%    5.552ms    200    27.76us    22.55us    33.15us    [CUDA m
0.12%    2.134ms    1    2.134ms    2.134ms    2.134ms    void th
```

**NVIDIA Provided**

**TAU**



Tuning and Analysis Utilities

**VampirTrace**



**3rd Party**

<https://developer.nvidia.com/performance-analysis-tools>

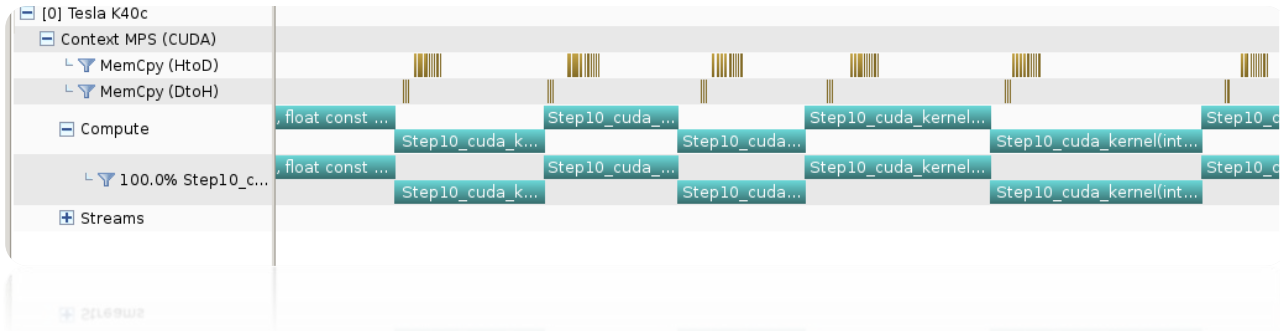
# NVPROF

## Command Line Profiler

- Compute time in each kernel
- Compute memory transfer time
- Collect metrics and events
- Support complex process hierarchy's
- Collect profiles for NVIDIA Visual Profiler
- No need to recompile

# NVIDIA's Visual Profiler (NVVP)

## Timeline



## Guided System

**1. CUDA Application Analysis**

**2. Performance-Critical Kernels**

**3. Compute, Bandwidth, or Latency Bound**

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "Step10\_cuda\_kernel" is most likely limited by compute.

**Perform Compute Analysis**

The most likely bottleneck to performance for this kernel is compute so you should first perform compute analysis to determine how it is limiting performance.

**Perform Latency Analysis**

**Perform Memory Bandwidth Analysis**

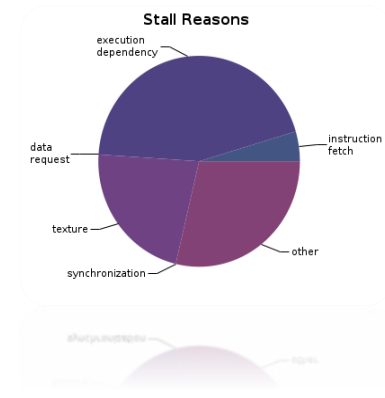
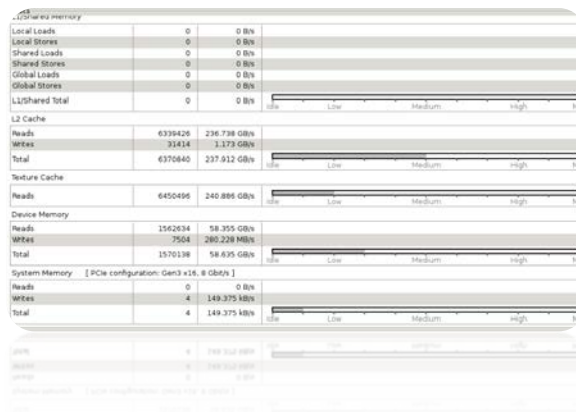
Instruction and memory latency and memory bandwidth are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

**Rerun Analysis**

If you modify the kernel you need to rerun your application to update this analysis.

go to the next step

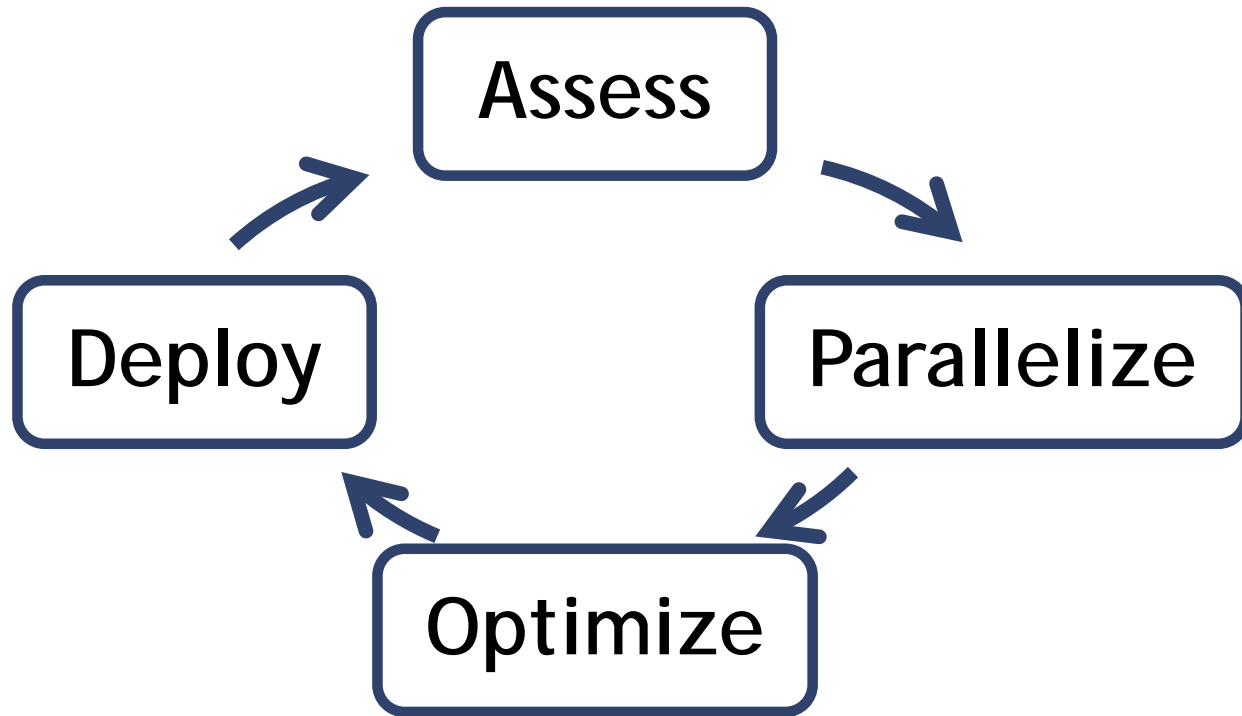
## Analysis



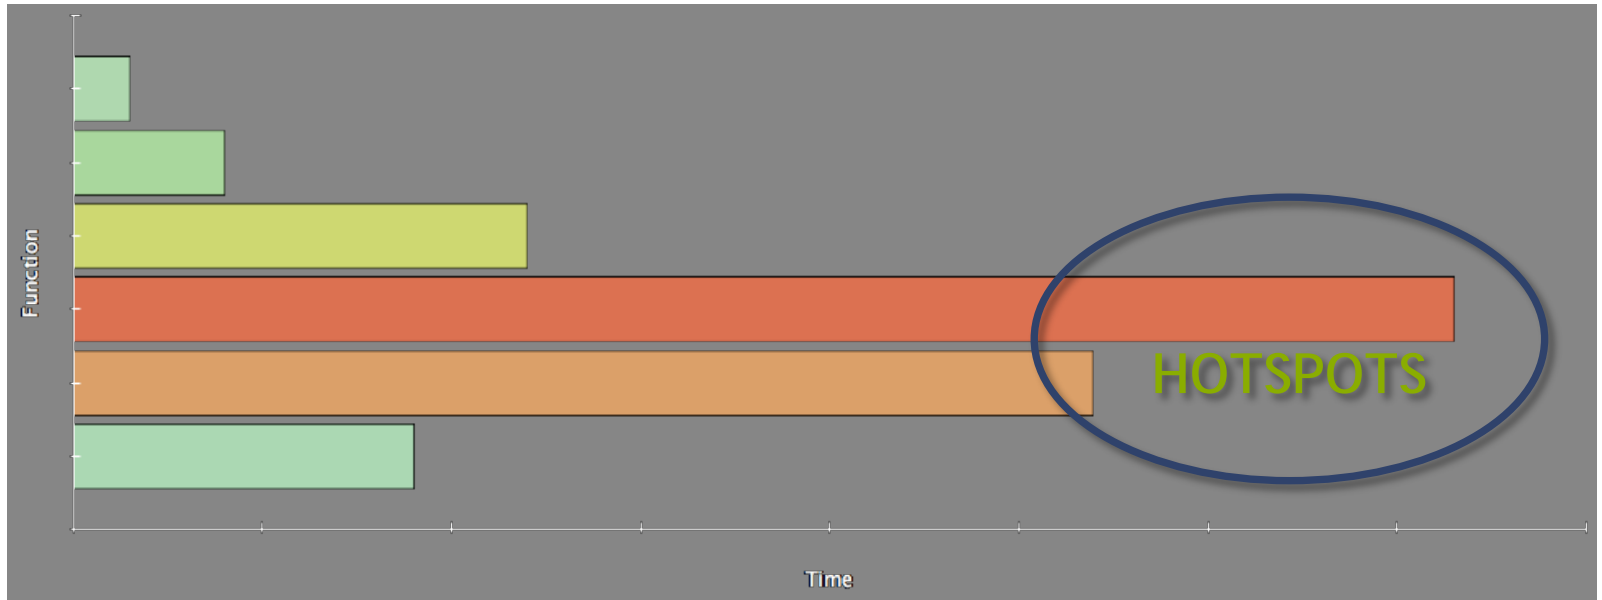
# Profiler Summary

- Many profile tools are available
- NVIDIA Provided
  - NVPROF: Command Line
  - NVVP: Visual profiler
  - NSIGHT: IDE (Visual Studio and Eclipse)
- 3<sup>rd</sup> Party
  - TAU
  - VAMPIR

# Optimization



# Assess



- Profile the code, find the hotspot(s)
- Focus your attention where it will give the most benefit

# Parallelize

Applications

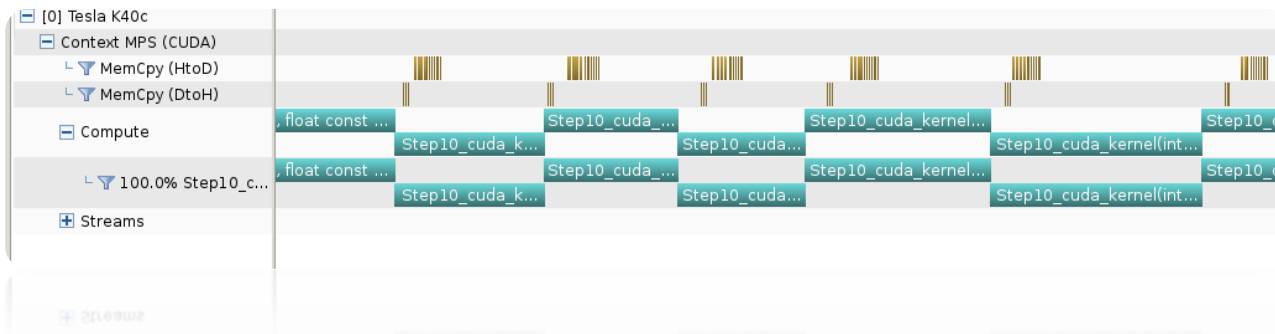
Libraries

Compiler  
Directives

Programming  
Languages

# Optimize

## Timeline



## Guided System

**1. CUDA Application Analysis**

**2. Performance-Critical Kernels**

**3. Compute, Bandwidth, or Latency Bound**

The first step in analyzing an individual kernel is to determine if the performance of the kernel is bounded by computation, memory bandwidth, or instruction/memory latency. The results at right indicate that the performance of kernel "Step10\_cuda\_kernel" is most likely limited by compute.

**Perform Compute Analysis**

The most likely bottleneck to performance for this kernel is compute so you should first perform compute analysis to determine how it is limiting performance.

**Perform Latency Analysis**

**Perform Memory Bandwidth Analysis**

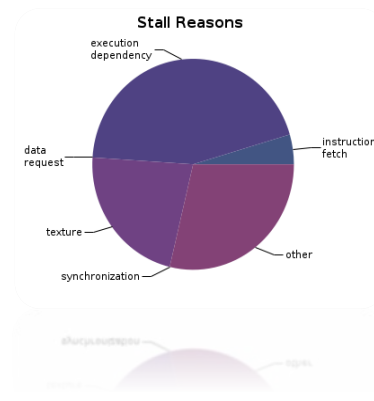
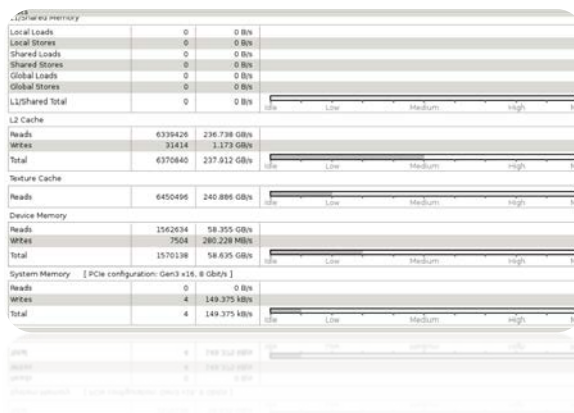
Instruction and memory latency and memory bandwidth are likely not the primary performance bottlenecks for this kernel, but you may still want to perform those analyses.

**Rerun Analysis**

If you modify the kernel you need to rerun your application to update this analysis.

go to the next slide to see the results of the analysis.

## Analysis





Obrigado!

# INTRODUÇÃO À PROGRAMAÇÃO PARA GPUs USANDO CUDA

---

Pedro Bruel

phrb@ime.usp.br

29 de Setembro de 2015



Instituto de Matemática e Estatística  
Universidade de São Paulo